

# CSC311, Fall 2022, Homework 2

Huipeng Lei, 1005908154

Due Monday, October 17, 2022

## Problem 1

a) The expected loss  $\mathbb{E}[\mathcal{J}(y, t)]$  for  $y = \text{Keep}$ :

$$\begin{aligned}\mathbb{E}[\mathcal{J}(y = \text{Keep}, t)] &= P(t = \text{NonSpam}) * \mathcal{J}(y = \text{Keep}, t = \text{NonSpam}) \\ &\quad + P(t = \text{Spam}) * \mathcal{J}(y = \text{Keep}, t = \text{Spam}) \\ &= 0 + 0.2 * 1 \\ &= 0.2\end{aligned}$$

The expected loss  $\mathbb{E}[\mathcal{J}(y, t)]$  for  $y = \text{Remove}$ :

$$\begin{aligned}\mathbb{E}[\mathcal{J}(y = \text{Remove}, t)] &= P(t = \text{NonSpam}) * \mathcal{J}(y = \text{Remove}, t = \text{NonSpam}) \\ &\quad + P(t = \text{Spam}) * \mathcal{J}(y = \text{Remove}, t = \text{Spam}) \\ &= 0.8 * 500 + 0 \\ &= 400\end{aligned}$$

b) Denote  $S = \text{Spam}$ ,  $NS = \text{NonSpam}$ ,  $K = \text{Keep}$ ,  $R = \text{Remove}$ .

Given the conditional probability  $P(t = S|x)$ , the Bayes optimal  $y_* \in \{\text{Keep}, \text{Remove}\}$  can be obtained by having  $y_* = \underset{y \in \{\text{Keep}, \text{Remove}\}}{\operatorname{argmin}} \mathbb{E}[\mathcal{J}(y, t)|x]$ . Where

When  $y = \text{Keep}$ :

$$\begin{aligned}\mathbb{E}[\mathcal{J}(y = K, t)|x] &= \mathcal{J}(y = K, t = NS) * P(t = NS|x) \\ &\quad + \mathcal{J}(y = K, t = S) * P(t = S|x) \\ &= 0 * P(t = NS|x) + 1 * P(t = S|x) \\ &= P(t = S|x)\end{aligned}$$

When  $y = \text{Remove}$ :

Note:  $P(t = S|x) = 1 - P(t = NS|x)$

$$\begin{aligned}\mathbb{E}[\mathcal{J}(y = R, t)|x] &= \mathcal{J}(y = R, t = NS) * P(t = NS|x) \\ &\quad + \mathcal{J}(y = R, t = S) * P(t = S|x) \\ &= 500 * P(t = NS|x) + 0 * P(t = S|x) \\ &= 500 * (1 - P(t = S|x))\end{aligned}$$

Since we want to minimize  $\mathbb{E}[\mathcal{J}(y, t)|x]$ , we can choose the Bayes optimal as

$$\begin{aligned}y_* &= \begin{cases} \text{Keep}, & P(t = S|x) \leq 500 * (1 - P(t = S|x)) \\ \text{Remove}, & P(t = S|x) > 500 * (1 - P(t = S|x)) \end{cases} \\ &= \begin{cases} \text{Keep}, & P(t = S|x) \leq \frac{500}{501} \\ \text{Remove}, & P(t = S|x) > \frac{500}{501} \end{cases}\end{aligned}$$

c) We know

$$P(x_1, x_2) = P((x_1, x_2)|t = S) * P(t = S) + P((x_1, x_2)|t = NS) * P(t = NS)$$

- $P(x_1 = 0, x_2 = 0) = 0.45 * 0.2 + 0.996 * 0.8 = 0.8868$
- $P(x_1 = 0, x_2 = 1) = 0.25 * 0.2 + 0.002 * 0.8 = 0.0516$
- $P(x_1 = 1, x_2 = 0) = 0.18 * 0.2 + 0.002 * 0.8 = 0.0376$
- $P(x_1 = 1, x_2 = 1) = 0.12 * 0.2 + 0 * 0.8 = 0.024$

From b), we only need to find what  $P(t = S|(x_1, x_2))$  is for different combinations of  $x_1$  and  $x_2$  and see which combination decides what the Bayes Optimal is. The formula is

$$P(t = S|(x_1, x_2)) = \frac{P((x_1, x_2)|t = S) * P(t = S)}{P(x_1, x_2)}$$

1. For  $x_1 = 0, x_2 = 0$ :

$$P(t = S|(x_1, x_2)) = \frac{0.45 * 0.2}{0.8868} = \frac{75}{739} = 0.101488... \leq \frac{500}{501}$$

2. For  $x_1 = 0, x_2 = 1$ :

$$P(t = S|(x_1, x_2)) = \frac{0.25 * 0.2}{0.0516} = \frac{125}{129} = 0.968992... \leq \frac{500}{501}$$

3. For  $x_1 = 1, x_2 = 0$ :

$$P(t = S|(x_1, x_2)) = \frac{0.18 * 0.2}{0.0376} = \frac{45}{47} = 0.9574468... \leq \frac{500}{501}$$

4. For  $x_1 = 1, x_2 = 1$ :

$$P(t = S|(x_1, x_2)) = \frac{0.12 * 0.2}{0.024} = 1 > \frac{500}{501}$$

Based on part b and the result we get, given features  $x_1, x_2$ , the Bayes Optimal decision for  $y_*$  would be

$$y_* = \begin{cases} \text{Keep}, & x_1 = 0, x_2 = 0 \vee x_1 = 0, x_2 = 1 \vee x_1 = 1, x_2 = 0 \\ \text{Remove}, & x_1 = 1, x_2 = 1 \end{cases}$$

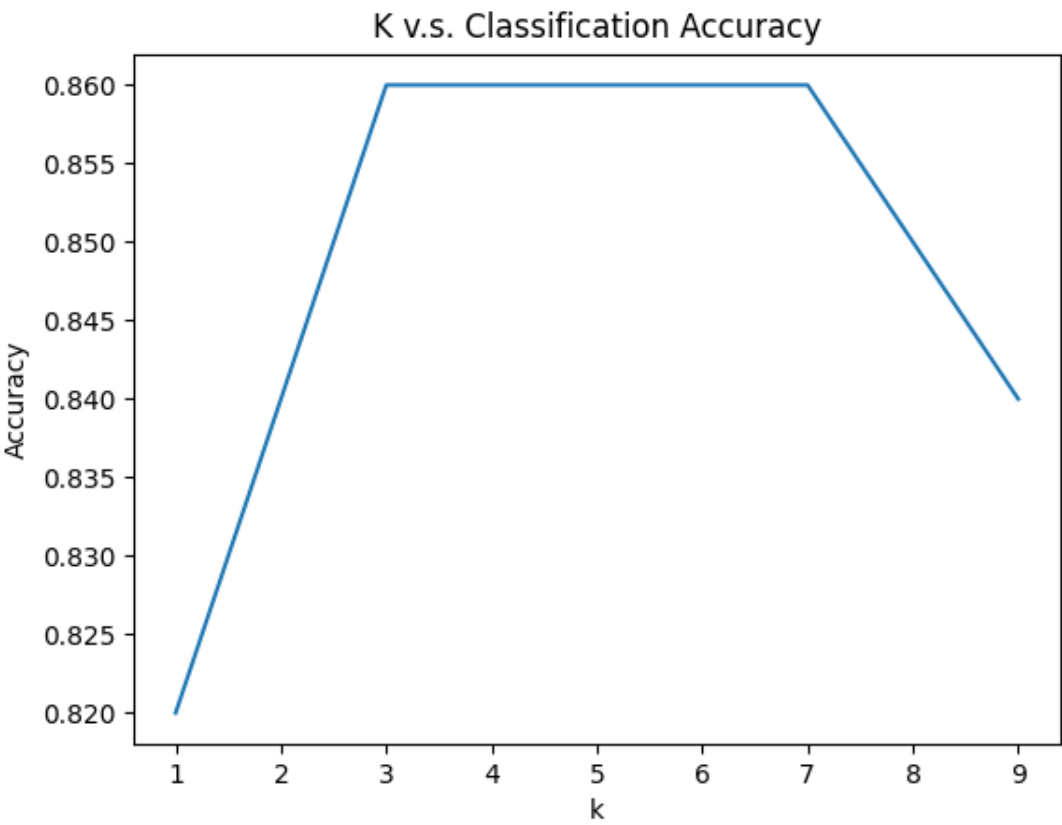
d) According to the Bayes decision rule from part c. The expected loss  $\mathbb{E}[\mathcal{J}(y_*, t)]$  is

$$\begin{aligned}
\mathbb{E}[\mathcal{J}(y_*, t)] &= 1 * P(y_* = K, t = S) + 500 * P(y_* = R, t = NS) \\
&= 1 * \left[ P(t = S | (x_1 = 0, x_2 = 0)) * P(x_1 = 0, x_2 = 0) \right. \\
&\quad + P(t = S | (x_1 = 0, x_2 = 1)) * P(x_1 = 0, x_2 = 1) \\
&\quad + P(t = S | (x_1 = 1, x_2 = 0)) * P(x_1 = 1, x_2 = 0) \left. \right] \\
&\quad + 500 * (1 - P(t = S | (x_1 = 1, x_2 = 1))) * P(x_1 = 1, x_2 = 1) \\
&= 1 * \left[ \frac{75}{739} * 0.8868 + \frac{125}{129} * 0.0516 + \frac{45}{47} * 0.0376 \right] + 500 * (1 - 1) * 0.024 \\
&= 0.176
\end{aligned}$$


## Problem 2

- a) From definition, a data-set is linearly separable if the training examples can be perfectly separated by a linear decision rule. However, if we want to visualize this data-set in Data Space, we see that the half-spaces are not convex. For instance, let's draw a line segment connecting two points  $(x_1 = -2, x_2 = -1)$  and  $(x_1 = 2, x_2 = 3)$  that lie in the positive half-space, we see that the line intersects  $(x_1 = 1, x_2 = 2)$ , which lies in the negative half-space. Therefore, the data-set is not linearly separable. Furthermore, there exists no linear line that separates this data-set perfectly.

### Problem 3



3.1 a)



b)

We chose  $k^* = 5$  since it has the highest validation accuracy for all  $k \in \{1, 3, 5, 7, 9\}$  as shown in part a. From the plot shown above, we see running K-Nearest-Neighbors on test inputs has higher accuracy than running KNN on validation inputs. The validation accuracy stays the same for all  $k^* - 2, k^*, k^* + 2$  at around 86%, but we see an increase of the test accuracy from 92% when  $k = k^* - 2$  to 94% for both  $k = k^*$  and  $k = k^* + 2$ . Moreover, as  $k$  changes, the test and validation accuracies change in the same direction (i.e. positively correlated).

3.2 b) The best hyper-parameter settings I've found is:

- Learning rate = 0.1
- Weight regularization = 0
- Number of iterations = 200
- Initialize weights =  $(M + 1) \times 1$  zero vector


Training set comes from `mnist_train`

```
{'learning_rate': 0.1, 'weight_regularization': 0.0, 'num_iterations': 200}
---Training---
Final Cross entropy: 0.0807561108521313
Classification Accuracy: 1.0
---Validation---
Final Cross entropy: 0.21857461101776998
Classification Accuracy: 0.88
---Test---
Final Cross entropy: 0.20737230681000285
Classification Accuracy: 0.92
```

Training set comes from `mnist_train_small`

```
{'learning_rate': 0.1, 'weight_regularization': 0.0, 'num_iterations': 200}
---Training---
Final Cross entropy: 0.011964472987210952
Classification Accuracy: 1.0
---Validation---
Final Cross entropy: 0.7496887121880731
Classification Accuracy: 0.7
---Test---
Final Cross entropy: 0.6701922932955665
Classification Accuracy: 0.78
```

- c) After running the code five times with the same hyperparameter settings, the result remains the same.



The cross-entropy decreases as the training progresses.



## Problem 4

- a) Let  $\mathcal{J}(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N a^{(i)} (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$ . We want to find  $\mathbf{w}^*$  that minimizes  $\mathcal{J}$  by setting the derivative of  $J$  w.r.t  $\mathbf{w}$  to 0. (i.e.  $\frac{\partial \mathcal{J}}{\partial \mathbf{w}} = 0$ )


Note:

$$\frac{\partial \mathbf{w}^T}{\partial \mathbf{w}} = 1$$

$$\frac{\partial \mathbf{w}^T \mathbf{w}}{\partial \mathbf{w}} = \frac{\partial \mathbf{w} \cdot \mathbf{w}}{\partial \mathbf{w}} = \frac{\partial \mathbf{w}}{\partial \mathbf{w}} * \mathbf{w} + \frac{\partial \mathbf{w}}{\partial \mathbf{w}} * \mathbf{w} = 2\mathbf{w}$$

Find  $\mathbf{w}^*$ :

$$\begin{aligned} \frac{\partial \mathcal{J}}{\partial \mathbf{w}} &= \sum_{i=1}^N a^{(i)} (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)}) (-\mathbf{x}^{(i)}) + \lambda \mathbf{w} \\ &= \sum_{i=1}^N (-\mathbf{x}^{(i)}) a^{(i)} (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)}) + \lambda \mathbf{w} \\ &= -\mathbf{X}^T \mathbf{A} (\mathbf{y} - \mathbf{X} \mathbf{w}) + \lambda \mathbf{w} \\ &= -\mathbf{X}^T \mathbf{A} \mathbf{y} + \mathbf{X}^T \mathbf{A} \mathbf{w} + \lambda \mathbf{w} \\ 0 &\stackrel{\text{set}}{=} -\mathbf{X}^T \mathbf{A} \mathbf{y} + (\mathbf{X}^T \mathbf{A} + \lambda \mathbf{I}) \mathbf{w} \\ (\mathbf{X}^T \mathbf{A} + \lambda \mathbf{I}) \mathbf{w} &= \mathbf{X}^T \mathbf{A} \mathbf{y} \\ \mathbf{w}^* &= (\mathbf{X}^T \mathbf{A} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{A} \mathbf{y} \end{aligned}$$



c)

- d) As  $\tau \rightarrow \inf$ , with  $\|x - x^{(i)}\|$  and  $\|x - x^{(j)}\|$  are fixed,  $a^{(i)} \rightarrow \frac{1}{N}$  because the terms in exponent tends to zero, i.e.  $\frac{\|x - x^{(i)}\|^2}{2\tau^2} \rightarrow 0$ . Therefore, the algorithm behaves almost the same as traditional  $l_2$  regularized linear regressions.

As  $\tau \rightarrow 0$ , the terms in exponent tends to infinity, which makes  $a^{(i)}$  super small. Since  $\tau$  acts as a bandwidth where it gives more weight to points that are closer to the test example, less to those that are further away. But since  $\tau \rightarrow 0$ , the weights for any points are super small, i.e. zero. In this case, the algorithm only focuses on the point exactly at the test example, which causes overfitting.

- e) **Advantage:** Locally weighted regression works well when there exists a non-linear relationship between the data and the targets, while traditional linear regression cannot be used to make prediction for a non-linear data-set.

**Disadvantage:** Higher computational cost than traditional linear regressions since it uses local fitting of data points and the model computes for each single data point.